

# Assignment 2

# Adding a user program to MavOS

## hello.c

```
#include "types.h"
#include "stat.h"
#include "user.h"

int main(int argc, char *argv[])
{
    printf("Hello, World\n");
    exit();
}
```

- The **user** directory is where all user space code and user applications live.
- Create your new program .c file in the user directory.



# Add the user program to the Makefile

```
119  UPROGS=\
120      $U/_cat\
121      $U/_echo\
122      $U/_forktest\
123      $U/_grep\
124      $U/_hello\
125      $U/_init\
126      $U/_kill\
127      $U/_ln\
128      $U/_ls\
129      $U/_mkdir\
130      $U/_rm\
131      $U/_sh\
132      $U/_stressfs\
133      $U/_usertests\
134      $U/_grind\
135      $U/_wc\
136      $U/_zombie\
137
```

- MacOS/Makefile contains the make rules to build our OS.
- In the UPROGS definition add a line for your new program.
- Note you need to prepend \_ to the filename.

# Steps to add a system call

- You will modify four files to add a new system call
  1. `syscall.h` - Reserve a system call number for your system call.
  2. `syscall.c` - Add an entry in the system call vector for your system call.
  3. `sysproc.c` - Add the system call function
  4. `usys.S` - Add a SYSCALL entry to allow the system to generate an entry point for your system call
  5. `user.h` - Add a function signature to the user space header file

# Reserving a system call number

```
1  #ifndef __SYSCALL_H_
2  #define __SYSCALL_H_
3
4  // System call numbers
5  #define SYS_fork    1
6  #define SYS_exit    2
7  #define SYS_wait    3
8  #define SYS_pipe    4
9  #define SYS_read    5
10 #define SYS_kill    6
11 #define SYS_exec    7
12 #define SYS_fstat    8
13 #define SYS_chdir    9
14 #define SYS_dup     10
15 #define SYS_getpid   11
16 #define SYS_sbrk     12
17 #define SYS_sleep    13
18 #define SYS_uptime   14
19 #define SYS_open     15
20 #define SYS_write    16
21 #define SYS_mknod    17
22 #define SYS_unlink   18
23 #define SYS_link     19
24 #define SYS_mkdir    20
25 #define SYS_close    21
26 #define SYS_mavup    22
27
28 #endif
29
```

- Claim the next integer for your system call by adding a new `#define` in `syscall.h`

# Add prototype for your system call

```
77 // Prototypes for the functions that handle system calls.
78 extern uint64 sys_fork(void);
79 extern uint64 sys_exit(void);
80 extern uint64 sys_wait(void);
81 extern uint64 sys_pipe(void);
82 extern uint64 sys_read(void);
83 extern uint64 sys_kill(void);
84 extern uint64 sys_exec(void);
85 extern uint64 sys_fstat(void);
86 extern uint64 sys_chdir(void);
87 extern uint64 sys_dup(void);
88 extern uint64 sys_getpid(void);
89 extern uint64 sys_sbrk(void);
90 extern uint64 sys_sleep(void);
91 extern uint64 sys_uptime(void);
92 extern uint64 sys_open(void);
93 extern uint64 sys_write(void);
94 extern uint64 sys_mknod(void);
95 extern uint64 sys_unlink(void);
96 extern uint64 sys_link(void);
97 extern uint64 sys_mkdir(void);
98 extern uint64 sys_close(void);
99 extern uint64 sys_mavup(void);
100
```

- Add a prototype for your system call function in `syscall.c`



# Add entry for your system call to the syscall vector

```
101 // An array mapping syscall numbers from syscall.h
102 // to the function that handles the system call.
103 static uint64 (*syscalls[])(void) = {
104     [SYS_fork] sys_fork,
105     [SYS_exit] sys_exit,
106     [SYS_wait] sys_wait,
107     [SYS_pipe] sys_pipe,
108     [SYS_read] sys_read,
109     [SYS_kill] sys_kill,
110     [SYS_exec] sys_exec,
111     [SYS_fstat] sys_fstat,
112     [SYS_chdir] sys_chdir,
113     [SYS_dup] sys_dup,
114     [SYS_getpid] sys_getpid,
115     [SYS_sbrk] sys_sbrk,
116     [SYS_sleep] sys_sleep,
117     [SYS_uptime] sys_uptime,
118     [SYS_open] sys_open,
119     [SYS_write] sys_write,
120     [SYS_mknod] sys_mknod,
121     [SYS_unlink] sys_unlink,
122     [SYS_link] sys_link,
123     [SYS_mkdir] sys_mkdir,
124     [SYS_close] sys_close,
125     [SYS_mavup] sys_mavup,
126 };
```

Add an entry for your system call function in the sys call vector

# Implement your system call function

```
1  #include "types.h"
2  #include "riscv.h"
3  #include "defs.h"
4  #include "param.h"
5  #include "memlayout.h"
6  #include "spinlock.h"
7  #include "process.h"
8
9  uint64 sys_mavup(void)
10 {
11     // mav up
12     return 0;
13 }
```

- In `sysproc.c` add the code for your system call function



# Now create the user space hooks

```
18  entry("fork");
19  entry("exit");
20  entry("wait");
21  entry("pipe");
22  entry("read");
23  entry("write");
24  entry("close");
25  entry("kill");
26  entry("exec");
27  entry("open");
28  entry("mknod");
29  entry("unlink");
30  entry("fstat");
31  entry("link");
32  entry("mkdir");
33  entry("chdir");
34  entry("dup");
35  entry("getpid");
36  entry("sbrk");
37  entry("sleep");
38  entry("uptime");
39  entry("mavup");
40
```

- In `user/usys.pl` add an entry for your system call.

# Add a function prototype in the user space header

```
3 // system calls
4 int fork(void);
5 int exit(int) __attribute__((noreturn));
6 int wait(int*);
7 int pipe(int*);
8 int write(int, const void*, int);
9 int read(int, void*, int);
10 int close(int);
11 int kill(int);
12 int exec(const char*, char**);
13 int open(const char*, int);
14 int mknod(const char*, short, short);
15 int unlink(const char*);
16 int fstat(int fd, struct stat*);
17 int link(const char*, const char*);
18 int mkdir(const char*);
19 int chdir(const char*);
20 int dup(int);
21 int getpid(void);
22 char* sbrk(int);
23 int sleep(int);
24 int uptime(void);
25 int mavup(void);
26 ..
```

- Add the function prototype to `user/user.h`

# Scheduler

```
451 void scheduler(void)
452 {
453     struct process_control_block *p;
454     struct cpu *c = mycpu();
455
456     c->proc = 0;
457     for(;;){
458         // Avoid deadlock by ensuring that devices can interrupt.
459         intr_on();
460
461         for(p = process_table; p < &process_table[NPROC]; p++)
462         {
463             acquire(&p->lock);
464             if(p->state == RUNNABLE)
465             {
466                 // Switch to chosen process. It is the process's job
467                 // to release its lock and then reacquire it
468                 // before jumping back to us.
469                 p->state = RUNNING;
470                 c->proc = p;
471                 swtch(&c->context, &p->context);
472
473                 // Process is done running for now.
474                 // It should have changed its p->state before coming back.
475                 c->proc = 0;
476             }
477             release(&p->lock);
478         }
479     }
480 }
```

- `process.c` contains the scheduler function `scheduler()`.
- The current scheduler implements round robin.
- Change the for loop to choose based on priority rules

# process control block

```
98 // Per-process state
99 ✓ struct process_control_block
100 {
101     struct spinlock lock;
102
103     // p->lock must be held when using these:
104     enum procstate state; // Process state
105     void *chan;           // If non-zero, sleeping on chan
106     int killed;           // If non-zero, have been killed
107     int xstate;           // Exit status to be returned to parent's wait
108     int pid;              // Process ID
109
110     // wait_lock must be held when using this:
111     struct process_control_block *parent; // Parent process
112
113     // these are private to the process, so p->lock need not be held.
114     uint64 kstack;         // Virtual address of kernel stack
115     uint64 sz;             // Size of process memory (bytes)
116     pagetable_t pagetable; // User page table
117     struct trapframe *trapframe; // data page for trampoline.S
118     struct context context; // swtch() here to run process
119     struct file *ofile[NOFILE]; // Open files
120     struct inode *cwd;        // Current directory
121     char name[16];           // Process name (debugging)
122 };
```

- The process control block is declared in `process.h`
- `allocproc(void)` allocates a new process and is a good spot to initialize the new process control block members.

# Common git message

```
⊗ @trevorbakker-uta → /workspaces/MavOS (main) $ git pull
remote: Enumerating objects: 22, done.
remote: Counting objects: 100% (22/22), done.
remote: Compressing objects: 100% (18/18), done.
remote: Total 20 (delta 10), reused 0 (delta 0), pack-reused 0 (from 0)
Unpacking objects: 100% (20/20), 553.55 KiB | 2.07 MiB/s, done.
From https://github.com/CSE3320-Fall-2024/MavOS
   ed74b27..b392026  main      -> origin/main
hint: You have divergent branches and need to specify how to reconcile them.
hint: You can do so by running one of the following commands sometime before
hint: your next pull:
hint:
hint:   git config pull.rebase false  # merge
hint:   git config pull.rebase true   # rebase
hint:   git config pull.ff only       # fast-forward only
hint:
hint: You can replace "git config" with "git config --global" to set a default
hint: preference for all repositories. You can also pass --rebase, --no-rebase,
hint: or --ff-only on the command line to override the configured default per
hint: invocation.
fatal: Need to specify how to reconcile divergent branches.
● @trevorbakker-uta → /workspaces/MavOS (main) $ git config pull.rebase false
```



Type: `git config pull.rebase false`